

✓ AI vs. Human-Generated Images Detection

Course: INFO-6147

Student: Harrison Kim, 1340629

Date: August 15, 2025

✓ 0. Environment Setup

```
1 !pip install -q torch torchvision kagglehub pandas tqdm matplotlib scikit-learn numpy
```

```
1 import torch
2
3 device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
4 print(f'Using device: {device}')
```

🔗 Using device: cuda

> 1. Dataset Selection:

- Choose a suitable image dataset for your project. You can consider any of the well-known datasets here [Datasets – Torchvision 0.17 documentation](#) for simplicity.
- Ensure that the dataset contains a reasonable number of classes and a sufficient number of images per class.
- If the dataset has very large number of images, you can use a subset (e.g., 1000 images per class if number of classes are 10 or less)
- If the dataset has more than 20 classes, you can use a subset of the classes (e.g., only use 10 classes)

[] ↳ 8 cells hidden

✓ 2. Data Preprocessing:

- Perform data preprocessing steps such as resizing images, normalizing pixel values, and splitting the dataset into training, validation, and test sets.
- Apply data augmentation techniques to increase the diversity of the training data.


✓ 2.1 Perform data preprocessing steps such as resizing images, normalizing pixel values, and splitting the dataset into training, validation, and test sets.

```
1 import torch
2 from torch.utils.data import Dataset
3 from PIL import Image
4
5 # Define ImageDataset class
6 class ImageDataset(Dataset):
7     def __init__(self, dataframe, transform=None):
8         self.dataframe = dataframe
9         self.transform = transform
10        self.image_paths = dataframe['image_path'].values
11        self.labels = dataframe['label'].values
12        self.label_to_idx = {label: idx for idx, label in enumerate(sorted(set(self.labels)))}
13
14    def __len__(self):
15        return len(self.image_paths)
16
```

```

17     def __getitem__(self, idx):
18         img_path = self.image_paths[idx]
19         label = self.labels[idx]
20         label_idx = self.label_to_idx[label]
21         try:
22             img = Image.open(img_path).convert('RGB')
23             if self.transform:
24                 img = self.transform(img)
25         except Exception as e:
26             print(f'Error loading image {img_path}:', e)
27             img = torch.zeros(3, 32, 32)
28         return img, label_idx
29
30 # 1. Split df into train, validation and test datasets
31 from sklearn.model_selection import train_test_split
32
33 temp = df[df['split'] == 'train'].reset_index(drop=True)
34 train_df, val_df = train_test_split(
35     temp,
36     test_size=0.2,
37     stratify=temp['label'],
38     random_state=42
39 )
40 test_df = df[df['split'] == 'test'].reset_index(drop=True)
41
42 # 2. Define transform (resize, normalize)
43 from torchvision import transforms
44 transform = transforms.Compose([
45     transforms.Resize((32, 32)),
46     transforms.ToTensor(),
47     transforms.Normalize(mean=[0.5, 0.5, 0.5], std=[0.5, 0.5, 0.5]),
48 ])
49
50 # 3. Create ImageDataset for each split
51 # train_dataset = ImageDataset(train_df, transform=transform)
52 val_dataset = ImageDataset(val_df, transform=transform)
53 test_dataset = ImageDataset(test_df, transform=transform)
54
55 print('Train set size:', len(train_df))
56 print('Validation set size:', len(val_dataset))
57 print('Test set size:', len(test_dataset))

```

 Train set size: 40000
 Validation set size: 10000
 Test set size: 10000

✓ 2.2 Apply data augmentation techniques to increase the diversity of the training data.

```

1 from torchvision import transforms
2
3 # Define data augmentation for training
4 train_transform = transforms.Compose([
5     transforms.Resize((32, 32)),
6     transforms.RandomHorizontalFlip(),
7     transforms.RandomRotation(10),
8     transforms.ToTensor(),
9     transforms.Normalize(mean=[0.5, 0.5, 0.5], std=[0.5, 0.5, 0.5]),
10 ])
11
12 # Re-create train_dataset with augmentation
13 train_dataset = ImageDataset(train_df, transform=train_transform)
14 print('Data augmentation applied to training dataset.')

```

 Data augmentation applied to training dataset.

✓ 3. Model Selection and Architecture:

- Select an appropriate deep learning architecture for image classification. You can start with a convolutional neural network (CNN).
- Define the architecture of your model, including the number of layers, activation functions, and any regularization techniques.

3.1 Select an appropriate deep learning architecture for image classification. You can start with a convolutional neural network (CNN).

- The CNN model uses Conv2d, ReLU, MaxPool2d, and Linear layers.

3.2 Define the architecture of your model, including the number of layers, activation functions, and any regularization techniques.

3.2.1 Define a custom CNN model with:

- **Conv2d:** This layer extracts key features from the input image, such as edges, textures, and shapes, which are essential for classification.
- **ReLU:** The activation function introduces non-linearity, enabling the model to learn more complex and abstract patterns beyond linear relationships.
- **MaxPool:** By reducing the spatial dimensions of the feature maps, this layer retains the most important information and helps prevent overfitting, while also improving computational efficiency.
- **Linear:** The fully connected layer combines all extracted features and produces the final output, determining the predicted class for each image.

```

1  import torch.nn as nn
2
3  class Net(nn.Module):
4      def __init__(self, num_classes):
5          super(Net, self).__init__()
6          self.features = nn.Sequential(
7              nn.Conv2d(3, 16, kernel_size=3, padding=1),
8              nn.ReLU(),
9              nn.MaxPool2d(2),
10             nn.Conv2d(16, 32, kernel_size=3, padding=1),
11             nn.ReLU(),
12             nn.MaxPool2d(2),
13             nn.Conv2d(32, 64, kernel_size=3, padding=1),
14             nn.ReLU(),
15             nn.MaxPool2d(2)
16         )
17         self.classifier = nn.Sequential(
18             nn.Flatten(),
19             nn.Linear(64 * 4 * 4, 128),
20             nn.ReLU(),
21             nn.Dropout(0.5),
22             nn.Linear(128, num_classes)
23         )
24
25     def forward(self, x):
26         x = self.features(x)
27         x = self.classifier(x)
28         return x
29
30 model = Net(num_classes=num_classes)
31
32 print("CNN model with dropout(0.5) initialized.")
33

```

 CNN model with dropout(0.5) initialized.

✓ 4. Model Training:

- Train your deep learning model using the training dataset.
- Monitor training progress, including loss and accuracy, and consider using early stopping to prevent overfitting.

```
1 # Initial hyperparameter values
2 BATCH_SIZE = 64
3 NUM_EPOCHS = 25
4 LR = 0.0001
```

✓ 4.1 Train your deep learning model using the training dataset.

✓ 4.1.1 Load dataset into dataloaders

```
1 from torch.utils.data import DataLoader
2
3 train_dataset = DataLoader(train_dataset, batch_size=BATCH_SIZE, shuffle=True)
4 val_dataset = DataLoader(val_dataset, batch_size=BATCH_SIZE, shuffle=False)
5 test_dataset = DataLoader(test_dataset, batch_size=BATCH_SIZE, shuffle=False)
```

✓ 4.1.2 Train custom CNN model

```
1 # Training loop for custom CNN model
2 import torch.optim as optim
3 from tqdm import tqdm
4 import copy
5
6 model = model.to(device)
7 criterion = nn.CrossEntropyLoss()
8 criterion = criterion.to(device)
9 optimizer = optim.Adam(model.parameters(), lr=LR)
10
11 # Lists to store metrics
12 cnn_train_loss_list = []
13 cnn_train_acc_list = []
14 cnn_val_loss_list = []
15 cnn_val_acc_list = []
16
17 best_val_acc = 0.0
18 counter = 0
19 for epoch in range(NUM_EPOCHS):
20     model.train()
21     running_loss = 0.0
22     correct = 0
23     total = 0
24     for images, labels in tqdm(train_dataset, desc=f"Epoch {epoch+1}/{NUM_EPOCHS}"):
25         images = images.to(device)
26         if not isinstance(labels, torch.Tensor):
27             labels = torch.tensor(labels)
28         labels = labels.to(device)
29         optimizer.zero_grad()
30         outputs = model(images)
31         loss = criterion(outputs, labels)
32         loss.backward()
33         optimizer.step()
34         running_loss += loss.item() * images.size(0)
35         _, predicted = torch.max(outputs, 1)
36         correct += (predicted == labels).sum().item()
37         total += labels.size(0)
38     train_loss = running_loss / total
39     train_acc = correct / total
40
```

```

41 # Store metrics
42 cnn_train_loss_list.append(train_loss)
43 cnn_train_acc_list.append(train_acc)
44
45 # Calculate and record validation loss
46 model.eval()
47 val_loss = 0.0
48 val_correct = 0
49 val_total = 0
50 with torch.no_grad():
51     for images, labels in val_dataset:
52         images = images.to(device)
53         if not isinstance(labels, torch.Tensor):
54             labels = torch.tensor(labels)
55         labels = labels.to(device)
56         outputs = model(images)
57         loss = criterion(outputs, labels)
58         val_loss += loss.item() * images.size(0)
59         _, predicted = torch.max(outputs, 1)
60         val_correct += (predicted == labels).sum().item()
61         val_total += labels.size(0)
62 val_loss = val_loss / val_total
63 val_acc = val_correct / val_total
64 cnn_val_loss_list.append(val_loss)
65 cnn_val_acc_list.append(val_acc)
66
67 if val_acc > best_val_acc:
68     best_val_acc = val_acc
69     counter = 0
70     best_model_state = copy.deepcopy(model.state_dict())
71 else:
72     counter += 1
73     if counter >= 5:
74         print("Early stopping triggered.")
75         break
76
77 print(f"Epoch {epoch+1}/{NUM_EPOCHS} | Train Loss: {train_loss:.4f} | Train Acc: {train_acc:.4f}")
78
79 print("\nTraining complete.")
80 print("cnn_train_loss_list:", cnn_train_loss_list)
81 print("cnn_train_acc_list:", cnn_train_acc_list)

```

```

Epoch 1/25: 100%|██████████| 625/625 [00:51<00:00, 12.17it/s]
Epoch 1/25 | Train Loss: 0.5261 | Train Acc: 0.7340
Epoch 2/25: 100%|██████████| 625/625 [00:45<00:00, 13.77it/s]
Epoch 2/25 | Train Loss: 0.4184 | Train Acc: 0.8108
Epoch 3/25: 100%|██████████| 625/625 [00:45<00:00, 13.66it/s]
Epoch 3/25 | Train Loss: 0.3809 | Train Acc: 0.8306
Epoch 4/25: 100%|██████████| 625/625 [00:45<00:00, 13.82it/s]
Epoch 4/25 | Train Loss: 0.3429 | Train Acc: 0.8543
Epoch 5/25: 100%|██████████| 625/625 [00:45<00:00, 13.83it/s]
Epoch 5/25 | Train Loss: 0.3133 | Train Acc: 0.8689
Epoch 6/25: 100%|██████████| 625/625 [00:45<00:00, 13.68it/s]
Epoch 6/25 | Train Loss: 0.2937 | Train Acc: 0.8777
Epoch 7/25: 100%|██████████| 625/625 [00:48<00:00, 12.79it/s]
Epoch 7/25 | Train Loss: 0.2796 | Train Acc: 0.8836
Epoch 8/25: 100%|██████████| 625/625 [00:47<00:00, 13.16it/s]
Epoch 8/25 | Train Loss: 0.2704 | Train Acc: 0.8885
Epoch 9/25: 100%|██████████| 625/625 [00:45<00:00, 13.88it/s]
Epoch 9/25 | Train Loss: 0.2607 | Train Acc: 0.8924
Epoch 10/25: 100%|██████████| 625/625 [00:45<00:00, 13.84it/s]
Epoch 10/25 | Train Loss: 0.2529 | Train Acc: 0.8968
Epoch 11/25: 100%|██████████| 625/625 [00:44<00:00, 14.14it/s]
Epoch 11/25 | Train Loss: 0.2461 | Train Acc: 0.9006
Epoch 12/25: 100%|██████████| 625/625 [00:45<00:00, 13.88it/s]
Epoch 12/25 | Train Loss: 0.2386 | Train Acc: 0.9032
Epoch 13/25: 100%|██████████| 625/625 [00:45<00:00, 13.68it/s]
Epoch 13/25 | Train Loss: 0.2340 | Train Acc: 0.9058
Epoch 14/25: 100%|██████████| 625/625 [00:46<00:00, 13.52it/s]
Epoch 14/25 | Train Loss: 0.2305 | Train Acc: 0.9075
Epoch 15/25: 100%|██████████| 625/625 [00:44<00:00, 13.93it/s]
Epoch 15/25 | Train Loss: 0.2271 | Train Acc: 0.9095
Epoch 16/25: 100%|██████████| 625/625 [00:45<00:00, 13.88it/s]

```

```

Epoch 16/25 | Train Loss: 0.2224 | Train Acc: 0.9131
Epoch 17/25: 100%|██████████| 625/625 [00:44<00:00, 14.01it/s]
Epoch 17/25 | Train Loss: 0.2152 | Train Acc: 0.9147
Epoch 18/25: 100%|██████████| 625/625 [00:44<00:00, 13.93it/s]
Epoch 18/25 | Train Loss: 0.2149 | Train Acc: 0.9171
Epoch 19/25: 100%|██████████| 625/625 [00:45<00:00, 13.80it/s]
Epoch 19/25 | Train Loss: 0.2100 | Train Acc: 0.9173
Epoch 20/25: 100%|██████████| 625/625 [00:47<00:00, 13.12it/s]
Epoch 20/25 | Train Loss: 0.2070 | Train Acc: 0.9183
Epoch 21/25: 100%|██████████| 625/625 [00:44<00:00, 14.00it/s]
Epoch 21/25 | Train Loss: 0.2036 | Train Acc: 0.9190
Epoch 22/25: 100%|██████████| 625/625 [00:44<00:00, 14.06it/s]
Epoch 22/25 | Train Loss: 0.1995 | Train Acc: 0.9213
Epoch 23/25: 100%|██████████| 625/625 [00:44<00:00, 14.05it/s]
Epoch 23/25 | Train Loss: 0.1995 | Train Acc: 0.9218
Epoch 24/25: 100%|██████████| 625/625 [00:43<00:00, 14.24it/s]
Epoch 24/25 | Train Loss: 0.1983 | Train Acc: 0.9220
Epoch 25/25: 100%|██████████| 625/625 [00:45<00:00, 13.87it/s]
Epoch 25/25 | Train Loss: 0.1943 | Train Acc: 0.9246

```

Training complete.

```

cnn_train_loss_list: [0.5260564805984497, 0.4183649462223053, 0.3809396066904068, 0.34290345828533175, 0.313298
cnn_train_acc_list: [0.734, 0.8108, 0.8306, 0.8543, 0.868875, 0.877725, 0.883625, 0.888475, 0.8924, 0.896825, 0.


```

✓ 4.2 Monitor training progress, including loss and accuracy, and consider using early stopping to prevent overfitting.

```

1 import matplotlib.pyplot as plt
2
3 plt.figure(figsize=(12,4))
4 plt.subplot(1,2,1)
5 plt.plot(cnn_train_loss_list, label='Train Loss')
6 plt.plot(cnn_val_loss_list, label='Val Loss')
7 plt.title('Loss (Train vs Validation)')
8 plt.xlabel('Epoch')
9 plt.ylabel('Loss')
10 plt.legend()
11
12 plt.subplot(1,2,2)
13 plt.plot(cnn_train_acc_list, label='Train Acc')
14 plt.plot(cnn_val_acc_list, label='Val Acc')
15 plt.title('Accuracy (Train vs Validation)')
16 plt.xlabel('Epoch')
17 plt.ylabel('Accuracy')
18 plt.legend()
19 plt.tight_layout()
20 plt.show()

```



✓ 5. Hyperparameter Tuning:

- Experiment with different hyperparameters (e.g., learning rate, batch size) to optimize the model's performance.
- Keep a record of the hyperparameters used and their impact on the model.

✓ 5.1 Experiment with different hyperparameters (e.g., learning rate, batch size) to optimize the model's performance.

```

1  # Hyperparameter search for learning rate and batch size
2  import copy
3  from tqdm import tqdm
4
5  search_learning_rates = [0.0005, 0.001, 0.005]
6  search_batch_sizes = [128, 192, 256]
7  num_epochs_hp = 5
8
9  results = []
10
11 for lr in search_learning_rates:
12     for batch_size in search_batch_sizes:
13         print(f"\nTraining with learning rate={lr}, batch size={batch_size}")
14         # Re-create dataloaders with new batch size
15         train_loader = DataLoader(train_dataset.dataset, batch_size=batch_size, shuffle=True)
16         val_loader = DataLoader(val_dataset.dataset, batch_size=batch_size, shuffle=False)
17
18         # Re-initialize model and optimizer
19         model = Net(num_classes=num_classes).to(device)
20         optimizer = optim.Adam(model.parameters(), lr=lr)
21
22         best_val_acc = 0.0
23         counter = 0
24         for epoch in range(num_epochs_hp):
25             model.train()
26             running_loss = 0.0
27             correct = 0
28             total = 0
29             for images, labels in tqdm(train_loader, desc=f"Epoch {epoch+1}/{num_epochs_hp} (train)"):
30                 images = images.to(device)
31                 if not isinstance(labels, torch.Tensor):
32                     labels = torch.tensor(labels)
33                 labels = labels.to(device)
34                 optimizer.zero_grad()
35                 outputs = model(images)
36                 loss = criterion(outputs, labels)
37                 loss.backward()
38                 optimizer.step()
39                 running_loss += loss.item() * images.size(0)
40                 _, predicted = torch.max(outputs, 1)
41                 correct += (predicted == labels).sum().item()
42                 total += labels.size(0)
43             train_loss = running_loss / total
44             train_acc = correct / total
45
46             # Validation
47             model.eval()
48             val_loss = 0.0
49             val_correct = 0
50             val_total = 0
51             with torch.no_grad():
52                 for images, labels in tqdm(val_dataset, desc="Validating"):
53                     images = images.to(device)
54                     if not isinstance(labels, torch.Tensor):
55                         labels = torch.tensor(labels)
56                     labels = labels.to(device)
57                     outputs = model(images)
58                     loss = criterion(outputs, labels)

```

```

59         val_loss += loss.item() * images.size(0)
60         _, predicted = torch.max(outputs, 1)
61         val_correct += (predicted == labels).sum().item()
62         val_total += labels.size(0)
63     val_loss = val_loss / val_total
64     val_acc = val_correct / val_total
65
66     if val_acc > best_val_acc:
67         best_val_acc = val_acc
68         counter = 0
69         best_model_state = copy.deepcopy(model.state_dict())
70     else:
71         counter += 1
72         if counter >= 3:
73             print("Early stopping triggered.")
74             break
75
76     print(f"\nBest validation accuracy: {best_val_acc:.4f}")
77     results.append({
78         'learning_rate': lr,
79         'batch_size': batch_size,
80         'best_val_acc': best_val_acc
81     })

```



Training with learning rate=0.0005, batch size=128
 Epoch 1/5 (train): 100% ██████████ 313/313 [00:43<00:00, 7.23it/s]
 Validating: 100% ██████████ 157/157 [00:09<00:00, 16.84it/s]
 Epoch 2/5 (train): 100% ██████████ 313/313 [00:43<00:00, 7.14it/s]
 Validating: 100% ██████████ 157/157 [00:09<00:00, 16.99it/s]
 Epoch 3/5 (train): 100% ██████████ 313/313 [00:49<00:00, 6.26it/s]
 Validating: 100% ██████████ 157/157 [00:09<00:00, 16.84it/s]
 Epoch 4/5 (train): 100% ██████████ 313/313 [00:44<00:00, 7.03it/s]
 Validating: 100% ██████████ 157/157 [00:08<00:00, 17.96it/s]
 Epoch 5/5 (train): 100% ██████████ 313/313 [00:43<00:00, 7.17it/s]
 Validating: 100% ██████████ 157/157 [00:09<00:00, 16.99it/s]
 Best validation accuracy: 0.9113

Training with learning rate=0.0005, batch size=192
 Epoch 1/5 (train): 100% ██████████ 209/209 [00:43<00:00, 4.84it/s]
 Validating: 100% ██████████ 157/157 [00:09<00:00, 16.80it/s]
 Epoch 2/5 (train): 100% ██████████ 209/209 [00:43<00:00, 4.79it/s]
 Validating: 100% ██████████ 157/157 [00:09<00:00, 17.05it/s]
 Epoch 3/5 (train): 100% ██████████ 209/209 [00:43<00:00, 4.77it/s]
 Validating: 100% ██████████ 157/157 [00:09<00:00, 17.24it/s]
 Epoch 4/5 (train): 100% ██████████ 209/209 [00:43<00:00, 4.85it/s]
 Validating: 100% ██████████ 157/157 [00:09<00:00, 16.80it/s]
 Epoch 5/5 (train): 100% ██████████ 209/209 [00:42<00:00, 4.94it/s]
 Validating: 100% ██████████ 157/157 [00:09<00:00, 17.41it/s]
 Best validation accuracy: 0.8983

Training with learning rate=0.0005, batch size=256
 Epoch 1/5 (train): 100% ██████████ 157/157 [00:42<00:00, 3.66it/s]
 Validating: 100% ██████████ 157/157 [00:08<00:00, 17.88it/s]
 Epoch 2/5 (train): 100% ██████████ 157/157 [00:43<00:00, 3.60it/s]
 Validating: 100% ██████████ 157/157 [00:08<00:00, 17.85it/s]
 Epoch 3/5 (train): 100% ██████████ 157/157 [00:42<00:00, 3.69it/s]
 Validating: 100% ██████████ 157/157 [00:09<00:00, 17.06it/s]
 Epoch 4/5 (train): 100% ██████████ 157/157 [00:42<00:00, 3.70it/s]
 Validating: 100% ██████████ 157/157 [00:09<00:00, 17.03it/s]
 Epoch 5/5 (train): 100% ██████████ 157/157 [00:43<00:00, 3.63it/s]
 Validating: 100% ██████████ 157/157 [00:08<00:00, 18.21it/s]
 Best validation accuracy: 0.9041

Training with learning rate=0.001, batch size=128
 Epoch 1/5 (train): 100% ██████████ 313/313 [00:44<00:00, 7.11it/s]
 Validating: 100% ██████████ 157/157 [00:09<00:00, 16.92it/s]
 Epoch 2/5 (train): 100% ██████████ 313/313 [00:43<00:00, 7.22it/s]
 Validating: 100% ██████████ 157/157 [00:08<00:00, 17.58it/s]
 Epoch 3/5 (train): 100% ██████████ 313/313 [00:43<00:00, 7.13it/s]
 Validating: 100% ██████████ 157/157 [00:09<00:00, 17.16it/s]
 Epoch 4/5 (train): 100% ██████████ 313/313 [00:43<00:00, 7.18it/s]
 Validating: 100% ██████████ 157/157 [00:09<00:00, 17.20it/s]


```
Epoch 5/5 (train): 100%|██████████| 313/313 [00:44<00:00, 7.07it/s]
Validating: 100%|██████████| 157/157 [00:08<00:00, 17.83it/s]
Best validation accuracy: 0.9261
```

```
Training with learning rate=0.001, batch size=192
Epoch 1/5 (train): 100%|██████████| 209/209 [00:43<00:00, 4.84it/s]
Validating: 100%|██████████| 157/157 [00:09<00:00, 16.27it/s]
Epoch 2/5 (train): 100%|██████████| 209/209 [00:43<00:00, 4.83it/s]
Validating: 100%|██████████| 157/157 [00:09<00:00, 17.06it/s]
```

✓ 5.2 Keep a record of the hyperparameters used and their impact on the model.

```
1 import pandas as pd
2
3 # Create a DataFrame to summarize hyperparameter search results
4 results_df = pd.DataFrame(results)
5 print("Hyperparameter summary table:")
6 print(results_df)
7
8 # Optionally, display the best configuration
9 best_row = results_df.loc[results_df['best_val_acc'].idxmax()]
10 print(f"\nBest configuration:\nLearning rate: {best_row['learning_rate']}, Batch size: {best_row['batch_size']}]
```

```
Hyperparameter summary table:
```

	learning_rate	batch_size	best_val_acc
0	0.0005	128	0.9113
1	0.0005	192	0.8983
2	0.0005	256	0.9041
3	0.0010	128	0.9261
4	0.0010	192	0.9208
5	0.0010	256	0.9219
6	0.0050	128	0.9121
7	0.0050	192	0.9212
8	0.0050	256	0.9209

```
Best configuration:
Learning rate: 0.001, Batch size: 128.0, Best Val Acc: 0.9261
```

✓ 6. Evaluation:

- Evaluate your trained model using the validation dataset to assess its performance.
- Calculate relevant metrics such as accuracy, precision, recall, and F1-score.
- Visualize the model's predictions and misclassifications.

✓ 6.1 Evaluate your trained model using the validation dataset to assess its performance.

```
1 # Evaluate CNN model on validation dataset
2 model.eval()
3 val_loss = 0.0
4 val_correct = 0
5 val_total = 0
6 cnn_true = []
7 cnn_pred = []
8 with torch.no_grad():
9     for images, labels in tqdm(val_dataset, desc="Validating"):
10         images = images.to(device)
11         if not isinstance(labels, torch.Tensor):
12             labels = torch.tensor(labels)
13         labels = labels.to(device)
14         outputs = model(images)
15         loss = criterion(outputs, labels)
16         val_loss += loss.item() * images.size(0)
17         _, predicted = torch.max(outputs, 1)
18         val_correct += (predicted == labels).sum().item()
19         val_total += labels.size(0)
```

```

20     cnn_true.extend(labels.cpu().numpy())
21     cnn_pred.extend(predicted.cpu().numpy())
22 val_loss = val_loss / val_total
23 val_acc = val_correct / val_total
24 print(f"\nCNN Validation Loss: {val_loss: 4f} | CNN Validation Accuracy: {val_acc: 4f}")
➦ Validating: 100%|██████████| 157/157 [00:09<00:00, 17.06it/s]CNN Validation Loss: 0.2834 | CNN Validation Accur

```

✓ 6.2 Calculate relevant metrics such as accuracy, precision, recall, and F1-score.

```

1 from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
2
3 # CNN metrics
4 cnn_acc = accuracy_score(cnn_true, cnn_pred)
5 cnn_prec = precision_score(cnn_true, cnn_pred, average='macro')
6 cnn_rec = recall_score(cnn_true, cnn_pred, average='macro')
7 cnn_f1 = f1_score(cnn_true, cnn_pred, average='macro')
8 print('CNN metrics:')
9 print(f'Accuracy: {cnn_acc:.4f}, Precision: {cnn_prec:.4f}, Recall: {cnn_rec:.4f}, F1-score: {cnn_f1:.4f}')

```

```

➦ CNN metrics:
Accuracy: 0.8758, Precision: 0.8920, Recall: 0.8758, F1-score: 0.8745

```

1 Start coding or [generate](#) with AI.

✓ 8. Final Model Testing:

- Test your final model on the held-out test dataset to assess its generalization to unseen data.

✓ 8.1 Test your final model on the held-out test dataset to assess its generalization to unseen data.

```

1 model.eval()
2 test_true = []
3 test_pred = []
4 with torch.no_grad():
5     for images, labels in test_dataset:
6         images = images.to(device)
7         if not isinstance(labels, torch.Tensor):
8             labels = torch.tensor(labels)
9         labels = labels.to(device)
10        outputs = model(images)
11        _, predicted = torch.max(outputs, 1)
12        test_true.extend(labels.cpu().numpy())
13        test_pred.extend(predicted.cpu().numpy())
14
15 test_acc = accuracy_score(test_true, test_pred)
16 test_prec = precision_score(test_true, test_pred, average='macro')
17 test_rec = recall_score(test_true, test_pred, average='macro')
18 test_f1 = f1_score(test_true, test_pred, average='macro')
19 print(f'Test Accuracy: {test_acc:.4f}, Precision: {test_prec:.4f}, Recall: {test_rec:.4f}, F1-score: {test_f1:

```

```

➦ Test Accuracy: 0.8759, Precision: 0.8919, Recall: 0.8759, F1-score: 0.8746

```

✓ 9. Documentation and Reporting:

- Create a project report summarizing your dataset, model architecture, training process, evaluation results, and insights gained.
- Include visualizations and explanations to make your findings clear.

- ✓ 9.1 Create a project report summarizing your dataset, model architecture, training process, evaluation results, and insights gained.

1 Start coding or [generate](#) with AI.

- ✓ 9.2 Include visualizations and explanations to make your findings clear.

1 Start coding or [generate](#) with AI.

✓ 10. Presentation:

- Prepare a brief presentation to showcase your project's key findings and outcomes.
- Share your experiences, challenges faced, and lessons learned during the project.

10.1 Prepare a brief presentation to showcase your project's key findings and outcomes.

10.2 Share your experiences, challenges faced, and lessons learned during the project.

✓ 11. Conclusion:

- Conclude your capstone project by summarizing your achievements and any future work or improvements that could be made to the model.
- Remember to maintain good coding practices and seek guidance or feedback from your instructor throughout the project.
- This capstone project will demonstrate your ability to apply deep learning techniques to real-world problems and showcase your skills to potential employers or collaborators.

11.1 Conclude your capstone project by summarizing your achievements and any future work or improvements that could be made to the model.

11.2 Remember to maintain good coding practices and seek guidance or feedback from your instructor throughout the project.

11.3 This capstone project will demonstrate your ability to apply deep learning techniques to real-world problems and showcase your skills to potential employers or collaborators.

